

NYCU Pattern Recognition, Homework 1

Deadline: March 23, 23:59

Part. 1, Coding (60%):

In this coding assignment, you need to implement linear regression by using only NumPy, then train your implemented model using **Gradient Descent** by the provided dataset and test the performance with testing data. Find the sample code and data on the GitHub page

https://github.com/NCTU-VRDL/CS_AT0828/tree/main/HW1

We suggest using the hyper-parameters below:

- Loss function: Mean Square Error
- Learning rate: $1e-4$
- Number of training iteration: 100

Please note that only NumPy can be used to implement your model, you will get no points by simply calling `sklearn.linear_model.LinearRegression`. Moreover, please train your regression model using Gradient Descent, not the closed-form solution.

1. (15%) Plot the learning curve of the training, you should find that loss decreases after a few iterations (x-axis=iteration, y-axis=loss, Matplotlib or other plot tools is available to use)
2. (15%) What's the Mean Square Error of your prediction and ground truth (prediction=model(x_test), ground truth=y_test)
3. (15%) What're the weights and intercepts of your linear model?
4. (10%) What's the difference between Gradient Descent, Mini-Batch Gradient Descent, and Stochastic Gradient Descent?
5. (5%) All your codes should follow the [PEP8 coding style](#) and with clear comments

Part. 2, Questions (40%):

1. (10%) Suppose that we have three colored boxes R (red), B (blue), and G (green). Box R contains 3 apples, 4 oranges, and 3 guavas, box B contains 2 apples, 0 orange, and 2 guavas, and box G contains 12 apples, 4 oranges, and 4 guavas. If a box is chosen at random with probabilities $p(R)=0.2$, $p(B)=0.4$, $p(G)=0.4$, and a piece of fruit is removed from the box (with equal probability of selecting any of the items in the box), then what is the probability of selecting guava? If we observe that the selected fruit is in fact an apple, what is the probability that it came from the blue box?
(3,4,3) (2,0,2) (12,4,4)
2. (15%) Consider two nonnegative numbers a and b , and show that, if $a \leq b$, then $a \leq (ab)^{1/2}$. Use this result to show that, if the decision regions of a two-class

classification problem are chosen to minimize the probability of misclassification, this probability will satisfy

$$p(\text{mistake}) \leq \int \{p(x, C_1) p(x, C_2)\}^{1/2} dx .$$

(Hint: Please refer to the textbook 1.5. Decision Theory)

3. (15%) Consider two variables x and y with joint distribution $p(x, y)$. Prove the following two results

$$\begin{aligned} E[x] &= E_y \left[E_x[x|y] \right] \\ \text{var}[x] &= E_y \left[\text{var}_x[x|y] \right] + \text{var}_y \left[E_x[x|y] \right] . \end{aligned}$$

Here $E_x[x|y]$ denotes the expectation of x under the conditional distribution $p(x|y)$, with a similar notation for the conditional variance.